

* Today's punchline:

Assume:

- ① $Y_i | \theta, \sigma^2 \sim N(\theta, \sigma^2)$
- ② $\theta | \sigma^2 \sim N(\mu_0, \sigma^2 / \kappa_0)$
- ③ $1/\sigma^2 \sim \text{gamma}(\frac{\nu_0}{2}, \frac{\nu_0}{2} \sigma_0^2)$

then the posterior

$$p(\theta, \sigma^2 | y_1, \dots, y_n)$$

$$= p(\theta | \sigma^2, y_1, \dots, y_n) p(\sigma^2 | y_1, \dots, y_n)$$

$$\downarrow \qquad \qquad \qquad \downarrow$$
$$\text{dnorm}(\theta; \mu_n, \tau_n) \text{dinv gamma}(\sigma^2; \frac{\nu_n}{2}, \frac{\nu_n}{2} \sigma_n^2)$$

"full cond'l posterior of θ "

Today's agenda:

- (1) sketch proof for $p(\sigma^2 | y_1, \dots, y_n)$
- (2) sample from ^{joint} posterior
- (3) sample from posterior predictive $p(\tilde{y} | y_1, \dots, y_n)$
↪ (time permitting)

Interpretation:

μ_0 : prior guess for θ

σ_0^2 : prior guess for σ^2

κ_0 : prior sample size for θ

ν_0 : prior sample size for ν_0